

00_Data_Wrangling

September 24, 2020

1 flats-in-cracow data wrangling

1.1 Imports

```
[1]: import matplotlib.pyplot as plt
import pandas as pd
import numpy as np

from collections import Counter
from IPython.display import display
from sklearn.impute import KNNImputer
from pylab import rcParams
from pathlib import Path
```

1.2 Setup

```
[2]: # Create directory for images
Path("img").mkdir(parents=True, exist_ok=True)

# Set default figure size
rcParams['figure.figsize'] = (4, 4)

# Tell pandas how to display floats
pd.options.display.float_format = "{:,.2f}".format
```

1.3 Goal

I scraped listings of properties for sale in Cracow. We would like to create a model to predict flat prices.

1.4 Data source

Data has been scraped from a website with listings. The data has undergone small transformations along the way. The goal of these transformations was to get the data into a usable state not to check it's validity.

1.5 Data loading

```
[3]: path = '../flats-data/raw_data.csv'
```

```
[4]: data = pd.read_csv(path, lineterminator='\n')
```

```
[5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 53110 entries, 0 to 53109
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  52975 non-null  object
1   City                  40097 non-null  object
2   District              29317 non-null  object
3   Amount                52924 non-null  float64
4   Currency              52924 non-null  object
5   Property              52612 non-null  object
6   Seller                52825 non-null  object
7   Area                  52683 non-null  float64
8   Rooms                 52071 non-null  float64
9   Bathrooms             34026 non-null  float64
10  Parking               22575 non-null  object
11  Garden                53110 non-null  bool
12  Balcony               53110 non-null  bool
13  Terrace               53110 non-null  bool
14  Floor                 53110 non-null  bool
15  New                   53110 non-null  bool
16  Estate                53110 non-null  bool
17  Townhouse             53110 non-null  bool
18  Apartment             53110 non-null  bool
19  Land                  53110 non-null  bool
20  Studio                53110 non-null  bool
21  Title                 52975 non-null  object
22  Description            46442 non-null  object
23  Link                  53110 non-null  object
dtypes: bool(10), float64(4), object(10)
memory usage: 6.2+ MB
```

First we sort the data in from newest to oldest, forcing rows with missing `Date` values to be last.

```
[6]: data = data.sort_values(by='Date',
                             ascending=False,
                             na_position='last',
                             ignore_index=True)
```

Next we assume that the `Title` column uniquely identifies a listing.

```
[7]: data = data.drop_duplicates(['Title'], keep='first')
```

After this the shape of the data is:

```
[8]: print(data.shape)
```

```
(9363, 24)
```

1.6 Data exploration

We check for missing values that we will have to deal with.

```
[9]: missing = data.isnull().sum(axis=0)
missing.name = 'Missing'
missing = missing.to_frame()
missing = missing[missing['Missing'] > 0]
missing.sort_values('Missing', ascending=False)
```

```
[9]:
```

	Missing
Parking	5867
District	3944
Bathrooms	3782
Description	1675
City	1565
Rooms	223
Area	121
Property	78
Seller	77
Amount	6
Currency	6
Date	1
Title	1

1.6.1 Check numeric columns

We see that we have 24 columns at our disposal. We inspect the numeric columns to see what we are dealing with. In the **Amount** column we note there is a property for sale that costs 1PLN, clearly a erroneous value. Next we note that the enourmous maximum in the **Amount** column. That is quite a lot of money and could be considered a potential outlier. The maximum and minimum of the **Area** column also indicate the existence of outliers. These values are clearly too large. The data will need to undergo a filtering process.

```
[10]: data.describe()
```

```
[10]:
```

	Amount	Area	Rooms	Bathrooms
count	9,357.00	9,242.00	9,140.00	5,581.00
mean	730,605.29	138.30	2.91	1.32
std	5,435,283.74	3,770.84	1.32	0.63

min	100.00	1.00	1.00	1.00
25%	395,000.00	42.00	2.00	1.00
50%	499,900.00	56.00	3.00	1.00
75%	720,000.00	80.00	4.00	1.00
max	521,290,000.00	320,000.00	6.00	4.00

1.6.2 Check binary columns

We inspect the data to see if binary columns are properly populated and check for imbalances.

```
[11]: binary = data.select_dtypes(bool).columns.to_list()

for col in binary:
    tmp = data[[col, 'Amount']]
    tmp = tmp.fillna('NaN')
    tmp = tmp.groupby(col, as_index=False)
    tmp = tmp.count()
    tmp = tmp.rename(columns={'Amount': 'Count'})
    tmp = tmp.sort_values('Count', ascending=False)
    tmp = tmp.reset_index(drop=True)
    display(tmp)
```

	Garden	Count
0	False	7564
1	True	1799

	Balcony	Count
0	False	6137
1	True	3226

	Terrace	Count
0	False	8275
1	True	1088

	Floor	Count
0	False	5767
1	True	3596

	New	Count
0	False	6370
1	True	2993

	Estate	Count
0	False	8021
1	True	1342

	Townhouse	Count
0	False	8539
1	True	824

	Apartment	Count
0	False	8016
1	True	1347

	Land	Count
0	False	7207
1	True	2156

	Studio	Count
0	False	8767
1	True	596

1.6.3 Check categorical columns

We inspect categorical columns to assert that they contain “valid” values. Most of these columns were generated by a script during the scraping and etl phase of the project.

```
[12]: categorical = data.select_dtypes('object').columns
categorical = categorical.to_list()
omit = ['Title', 'Link', 'Description', 'Date']

for col in categorical:
    if col not in omit:
        tmp = data[['Amount', col]].copy()
        tmp = tmp.fillna('NaN')
        tmp = tmp.groupby(col, as_index=False)
        tmp = tmp.count()
        tmp = tmp.rename(columns={'Amount': 'Count'})
        tmp = tmp.sort_values('Count', ascending=False)
        tmp = tmp.reset_index(drop=True)
        display(tmp)
```

	City	Count
0	kraków	7798
1	NaN	1565

	District	Count
0	NaN	3944
1	krowodrza	724
2	stare miasto	623

3	podgorze	591
4	nowa huta	414
5	debniki	399
6	bronowice	388
7	pradnik bialy	360
8	biezanow	291
9	pradnik czerwony	284
10	grzegorzki	269
11	czyzyny	207
12	mistrzejowice	169
13	lagiewniki	138
14	zwierzyniec	129
15	podgorze duchackie	120
16	bienczyce	109
17	swoszowice	96
18	prokocim	59
19	borek falecki	31
20	wzgorza krzeslawickie	18

	Currency	Count
0	pln	9357
1	NaN	6

	Property	Count
0	flat	8055
1	house	1230
2	NaN	78

	Seller	Count
0	realtor	8569
1	owner	717
2	NaN	77

	Parking	Count
0	NaN	5867
1	street	1418
2	garage	1331
3	no parking	573
4	covered	174

1.6.4 Check text columns

We search for keywords in the data.

```
[13]: # text = data[data['Description'].isna() == False].copy()
# text = text['Description'].to_list()
# text = ' '.join(text)
# text = text.split(' ')
# text = [x for x in text if x.isalpha()]
# text = [x for x in text if len(x) > 3]
```

```
[14]: # for i in range(5, len(text)-5):
#     if 'piętro' in text[i]:
#         s = text[i-5:i+5]
#         s = ' '.join(s)
#         print(s)
```

1.7 Data cleaning

We assume that if we know the district, the City is **kraków**.

```
[15]: mask = (data['City'].isna() == True) & (data['District'].isna() == False)
data.loc[mask, 'City'] = 'kraków'
```

We extract more Parking information from the property description.

```
[16]: def extract_parking(x):
    if ('garaż' in x or 'garaz' in x or 'parking' in x) and 'podziemny' in x:
        return 'covered'
    elif ('garaż' in x or 'garaz' in x) and 'podziemny' not in x:
        return 'garage'
    elif 'parking' in x and 'podziemny' not in x:
        return 'street'
    else:
        return 'no parking'
```

```
[17]: mask = (data['Parking'].isna() == True) & (data['Description'].isna() == False)
data.loc[mask, ['Parking', 'Description']] = data.loc[mask, 'Description'].
    ↪apply(extract_parking)
```

```
[18]: mask = data['Parking'].isna() == True
data.loc[mask, 'Parking'] = 'no parking'
```

We confirm that we have dealt with all the NaNs in the Parking column.

```
[19]: print(data['Parking'].isna().sum())
```

0

1.7.1 Filtering

Next we filter the data according to these rules:

```
[20]: data = data[data['City'] == 'kraków']
data = data[data['Currency'] == 'pln']
data = data[data['Property'] == 'flat']
data = data[(data['Amount'] >= data['Amount'].quantile(0.025))]
data = data[(data['Amount'] <= data['Amount'].quantile(0.975))]
data = data[(data['Area'] >= data['Area'].quantile(0.01))]
data = data[(data['Area'] <= data['Area'].quantile(0.99))]
data = data[data['District'] != 'unknown']
data = data[data['District'].isna() == False]
data = data[data['Seller'].isna() == False]
data = data[data['Description'].isna() == False]
```

```
[21]: data = data.reset_index(drop=True)
```

```
[22]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4102 entries, 0 to 4101
Data columns (total 24 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Date            4102 non-null  object
1   City            4102 non-null  object
2   District        4102 non-null  object
3   Amount          4102 non-null  float64
4   Currency        4102 non-null  object
5   Property        4102 non-null  object
6   Seller          4102 non-null  object
7   Area            4102 non-null  float64
8   Rooms           4047 non-null  float64
9   Bathrooms       1971 non-null  float64
10  Parking         4102 non-null  object
11  Garden          4102 non-null  bool
12  Balcony         4102 non-null  bool
13  Terrace         4102 non-null  bool
14  Floor           4102 non-null  bool
15  New             4102 non-null  bool
16  Estate          4102 non-null  bool
17  Townhouse       4102 non-null  bool
18  Apartment       4102 non-null  bool
19  Land            4102 non-null  bool
20  Studio          4102 non-null  bool
21  Title           4102 non-null  object
22  Description      4102 non-null  object
23  Link            4102 non-null  object
dtypes: bool(10), float64(4), object(10)
memory usage: 488.8+ KB
```


1.7.2 Impute missing values

The next step is to fill in missing values for numeric columns `Amount`, `Area`, `Rooms` and `Bathrooms`. We use the `KNNImputer` to accomplish this.

```
[23]: numeric = list(data.select_dtypes('number').columns)

[24]: mask = (data['Bathrooms'].isna() == True | data['Rooms'].isna())
      missing = data[numeric]

      imputer = KNNImputer(n_neighbors=5)
      imputer.fit(missing)

      missing = imputer.transform(missing)
      missing = pd.DataFrame(missing, columns=numeric)

      for col in numeric:
          data[col] = missing[col]

      for col in numeric:
          data[col] = data[col].apply(lambda x: round(x))

[25]: print(data.shape)
```

```
(4102, 24)
```

1.8 Save data

Verify that there are no NaNs in data.

```
[26]: data.isnull().sum().sum()
```

```
[26]: 0
```

Remove columns that will not be used further.

```
[27]: data = data.drop(['Title',
                       'Description',
                       'Link',
                       'Property',
                       'City',
                       'Currency',
                       'Date'], axis=1)
```

Take a last peek at the data.

```
[28]: data.head()
```

```
[28]:
```

	District	Amount	Seller	Area	Rooms	Bathrooms	Parking	Garden	\
0	biezanow	439082	realtor	56	3	1	covered	True	
1	podgorze	845000	realtor	132	5	2	no parking	False	
2	podgorze	360000	realtor	41	2	1	no parking	False	
3	krowodrza	1190000	realtor	81	3	1	no parking	True	
4	debniki	990000	realtor	93	4	2	street	False	

	Balcony	Terrace	Floor	New	Estate	Townhouse	Apartment	Land	Studio
0	True	False	True	True	False	False	True	True	False
1	False	True	False	False	False	False	True	False	False
2	True	False	False	False	False	False	False	False	False
3	False	True	True	True	False	False	True	True	False
4	False	False	False	False	False	False	False	False	False

```
[29]: data.describe()
```

```
[29]:
```

	Amount	Area	Rooms	Bathrooms
count	4,102.00	4,102.00	4,102.00	4,102.00
mean	532,119.80	55.62	2.60	1.10
std	221,445.75	20.06	0.99	0.32
min	210,000.00	23.00	1.00	1.00
25%	390,000.00	41.00	2.00	1.00
50%	470,000.00	53.00	3.00	1.00
75%	610,000.00	66.00	3.00	1.00
max	1,525,000.00	134.00	6.00	4.00

Save it for further analysis.

```
[30]: data.to_csv('../flats-data/cleaned_data.csv', index=False)
```