

02_Model

September 24, 2020

1 flats-in-cracow machine learning

1.1 Imports

```
[1]: from datetime import datetime
from distutils.dir_util import copy_tree
from pathlib import Path

import joblib
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

from matplotlib.ticker import MaxNLocator
from pylab import rcParams
from sklearn.compose import ColumnTransformer, TransformedTargetRegressor
from sklearn.dummy import DummyRegressor
from sklearn.ensemble import (GradientBoostingRegressor, RandomForestRegressor,
                              VotingRegressor)
from sklearn.impute import KNNImputer
from sklearn.metrics import (mean_absolute_error, mean_squared_error,
                              mean_squared_log_error)
from sklearn.model_selection import GridSearchCV, train_test_split, KFold
from sklearn.neural_network import MLPRegressor
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import MinMaxScaler, OneHotEncoder
```

1.2 Setup

```
[2]: # Create directory for images
Path("img").mkdir(parents=True, exist_ok=True)

# Set default figure size
rcParams['figure.figsize'] = (4, 4)

# Tell pandas how to display floats
pd.options.display.float_format = "{:,.2f}".format
```

1.3 Data loading

```
[3]: path = '../flats-data/cleaned_data.csv'
```

```
[4]: data = pd.read_csv(path, lineterminator='\n')
```

```
[5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4102 entries, 0 to 4101
Data columns (total 17 columns):
#   Column          Non-Null Count  Dtype
---  -
0   District        4102 non-null   object
1   Amount          4102 non-null   int64
2   Seller          4102 non-null   object
3   Area            4102 non-null   int64
4   Rooms           4102 non-null   int64
5   Bathrooms       4102 non-null   int64
6   Parking         4102 non-null   object
7   Garden          4102 non-null   bool
8   Balcony         4102 non-null   bool
9   Terrace         4102 non-null   bool
10  Floor           4102 non-null   bool
11  New             4102 non-null   bool
12  Estate          4102 non-null   bool
13  Townhouse       4102 non-null   bool
14  Apartment       4102 non-null   bool
15  Land            4102 non-null   bool
16  Studio          4102 non-null   bool
dtypes: bool(10), int64(4), object(3)
memory usage: 264.5+ KB
```

```
[6]: data.head()
```

```
[6]:
```

	District	Amount	Seller	Area	Rooms	Bathrooms	Parking	Garden	\
0	biezanow	439082	realtor	56	3	1	covered	True	
1	podgorze	845000	realtor	132	5	2	no parking	False	
2	podgorze	360000	realtor	41	2	1	no parking	False	
3	krowodrza	1190000	realtor	81	3	1	no parking	True	
4	debniki	990000	realtor	93	4	2	street	False	

	Balcony	Terrace	Floor	New	Estate	Townhouse	Apartment	Land	Studio
0	True	False	True	True	False	False	True	True	False
1	False	True	False	False	False	False	True	False	False
2	True	False	False	False	False	False	False	False	False
3	False	True	True	True	False	False	True	True	False

4 False False False False False False False False

1.4 Feature engineering

The next step is to engineer features. We add columns describing the `Total Rooms` in the property, ratio of `Area` to `Rooms` and so on.

```
[7]: data['Log Area'] = np.round(np.log(data['Area']), 2)
data['Bool Sum'] = data.select_dtypes(bool).sum(axis=1)
# Avoid division by zero
data['Area to Bool Sum'] = round(data['Area'] / (data.select_dtypes(bool).
    ↳sum(axis=1) + 1), 2)
data['Rooms to Bool Sum'] = round(data['Rooms'] / (data.select_dtypes(bool).
    ↳sum(axis=1) + 1), 2)
data['Rooms to Bathrooms'] = round(data['Rooms'] / data['Bathrooms'], 2)
data['Total Rooms'] = round(data['Rooms'] + data['Bathrooms'], 2)
data['Area to Rooms'] = round(data['Area'] / data['Rooms'], 2)
data['Area to Bathrooms'] = round(data['Area'] / data['Rooms'], 2)
data['Area to Total Rooms'] = round(data['Area'] / data['Total Rooms'], 2)
```

1.5 Data split

We decide to use 80% of the data to train the model and 20% to check performance. We make sure to remove the `Amount` column from the training data since this is our target and remove duplicates before training.

```
[8]: print(len(data))
data = data.drop_duplicates()
print(len(data))
```

4102

3586

```
[9]: X = data.drop(['Amount'], axis=1)
y = data['Amount']

split = train_test_split(X, y, train_size=.8,
                        random_state=123)

X_train, X_test, y_train, y_test = split
```

1.6 Models

Next step is to create the models and associated pipelines. We apply one hot encoding to categorical features and use the `ColumnTransformer` parameter `passthrough` to allow the rest of the columns to remain unchanged.

```
[10]: categorical = list(X.select_dtypes('object').columns)
      continuous = list(X.select_dtypes('int64'))
      continuous += list(X.select_dtypes('float64'))
```

1.6.1 Baseline model

For comparison purposes we create a model to give base predictions.

```
[11]: dmr = DummyRegressor()

      dmr_ohe = Pipeline(
          steps=[('onehot', OneHotEncoder(handle_unknown='ignore'))])

      dmr = Pipeline(steps = [('preprocessor', dmr_ohe),
                              ('regressor', dmr)])

      # dmr.fit(X_train, y_train)
```

1.6.2 Multi-layer Perceptron

For the neural network we apply the `MinMaxScaler` so that the continuous columns have values in $[0, 1]$ and then we apply `OneHotEncoder` to the categorical columns.

```
[12]: mlp = MLPRegressor(hidden_layer_sizes=(100, 100, 100),
                          max_iter=2*10**4,
                          random_state=123)

      mlp_ohe = Pipeline(
          steps=[('onehot', OneHotEncoder(handle_unknown='ignore'))])

      mlp_scale = Pipeline(
          steps=[('scale', MinMaxScaler())])

      mlp_pre = ColumnTransformer(
          transformers = [
              ('scale', mlp_scale, continuous),
              ('cat', mlp_ohe, categorical),
          ],
          remainder='passthrough'
      )

      mlp_trans = TransformedTargetRegressor(regressor=mlp,
                                              transformer=MinMaxScaler())
```

```
mlp = Pipeline(steps = [('preprocessor', mlp_pre),
                        ('transformer', mlp_trans)])

# mlp.fit(X_train, y_train)
```

1.6.3 Gradient Boosting Regressor

For the gradient booster we only apply `OneHotEncoder` to the categorical columns.

```
[13]: gbr = GradientBoostingRegressor(random_state=123)

gbr_ohe = Pipeline(
    steps=[('onehot', OneHotEncoder(handle_unknown='ignore'))])

gbr_pre = ColumnTransformer(
    transformers = [
        ('cat', gbr_ohe, categorical)
    ],
    remainder='passthrough'
)

gbr = Pipeline(steps = [('preprocessor', gbr_pre),
                        ('regressor', gbr)])

# gbr.fit(X_train, y_train)
```

1.7 Parameter tuning

We set up the training process to conduct basic parameter tuning and cross validation.

```
[14]: kf = KFold(n_splits=5, random_state=123, shuffle=True)

[15]: gbr_grid = {'regressor__max_depth': [5, 10, 15],
                  'regressor__n_estimators': [50, 100, 200, 300],
                  'regressor__min_samples_split': [2, 4],
                  'regressor__min_samples_leaf': [2, 4],
                  'regressor__max_features': ['auto']}

gbr_gs = GridSearchCV(estimator=gbr,
                      param_grid=gbr_grid,
                      cv=kf,
                      n_jobs=8,
                      scoring='neg_root_mean_squared_error',
                      verbose=2)
```

```
[16]: layers = [(100, 100, 100),
               (150, 200, 150),
               (200, 400, 200)]

mlp_grid = {'transformer__regressor__activation': ['relu'],
            'transformer__regressor__solver': ['adam'],
            'transformer__regressor__learning_rate': ['adaptive'],
            'transformer__regressor__learning_rate_init': [0.01, 0.001, 0.0001],
            'transformer__regressor__hidden_layer_sizes': layers}

mlp_gs = GridSearchCV(estimator=mlp,
                      param_grid=mlp_grid,
                      cv=kf,
                      n_jobs=8,
                      scoring='neg_root_mean_squared_error',
                      verbose=2)
```

1.8 Training

```
[17]: dmr.fit(X_train, y_train)
```

```
[17]: Pipeline(steps=[('preprocessor',
                       Pipeline(steps=[('onehot',
                                         OneHotEncoder(handle_unknown='ignore'))])),
                      ('regressor', DummyRegressor())])
```

```
[18]: mlp = mlp_gs.fit(X_train, y_train).best_estimator_
mlp
```

Fitting 5 folds for each of 9 candidates, totalling 45 fits

[Parallel(n_jobs=8)]: Using backend LokyBackend with 8 concurrent workers.

[Parallel(n_jobs=8)]: Done 25 tasks | elapsed: 18.6s

[Parallel(n_jobs=8)]: Done 45 out of 45 | elapsed: 47.5s finished

```
[18]: Pipeline(steps=[('preprocessor',
                      ColumnTransformer(remainder='passthrough',
                                      transformers=[('scale',
                                                    Pipeline(steps=[('scale',
                                                                      MinMaxScaler()))]),
                                                    ('Area', 'Rooms', 'Bathrooms',
                                                                 'Bool Sum', 'Total Rooms',
                                                                 'Log Area',
                                                                 'Area to Bool Sum',
                                                                 'Rooms to Bool Sum',
                                                                 'Rooms to Bathrooms',
                                                                 'Area to Rooms',
```

```

        'Area to Bathrooms',
        'Area to Total Rooms']],
('cat',
 Pipeline(steps=[('onehot',
                    OneHotEncoder(handle_unknown='ignore'))]),
            ['District', 'Seller',
             'Parking']]))),
        ('transformer',
 TransformedTargetRegressor(regressor=MLPRegressor(hidden_layer_sizes=(150,
        200,
        150),
        learning_rate='adaptive',
        learning_rate_init=0.01,
        max_iter=20000,
        random_state=123),
                    transformer=MinMaxScaler()))))

```

CV RMSE score for MLPRegressor:

```
[19]: print(round(abs(mlp_gs.best_score_)))
```

118596

```
[20]: gbr = gbr_gs.fit(X_train, y_train).best_estimator_
      gbr
```

Fitting 5 folds for each of 48 candidates, totalling 240 fits

```

[Parallel(n_jobs=8)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=8)]: Done 25 tasks      | elapsed:    7.8s
[Parallel(n_jobs=8)]: Done 146 tasks    | elapsed:   1.1min
[Parallel(n_jobs=8)]: Done 240 out of 240 | elapsed:   2.6min finished

```

```

[20]: Pipeline(steps=[('preprocessor',
                        ColumnTransformer(remainder='passthrough',
                                           transformers=[('cat',
                                                            Pipeline(steps=[('onehot',
                                                                                   OneHotEncoder(handle_unknown='ignore'))]),
                                                                                   ['District', 'Seller',
                                                                                   'Parking']]))),
                        ('regressor',
 GradientBoostingRegressor(max_depth=5, max_features='auto',
                             min_samples_leaf=2,
                             random_state=123))])

```

CV RMSE score for GradientBoostingRegressor:

```
[21]: print(round(abs(gbr_gs.best_score_)))
```

121466

1.8.1 Voting Regressor

We create a `VotingRegressor` with uniform weights to be able to combine predictions of our models.

```
[22]: vote = VotingRegressor(estimators=[['mlp', mlp], ['gbr', gbr]], n_jobs=8)
vote = vote.fit(X_train, y_train)
```

1.9 Model performance

We obtain predictions for the testing set and compare RMSE, MAE and MSLE scores of our models.

```
[23]: def get_scores(regressor, X_test, y_true, verb=True):
    """
    Obtain RMSE, MAE and MSLE for test set.
    """

    y_pred = regressor.predict(X_test)

    rmse = mean_squared_error(y_pred=y_pred,
                              y_true=y_test,
                              squared=False)

    mae = mean_absolute_error(y_pred=y_pred,
                              y_true=y_test)

    msle = mean_squared_log_error(y_pred=y_pred,
                                   y_true=y_test)

    if verb:

        print(f'RMSE: {rmse:10.2f}')
        print(f'MAE: {mae:10.2f}')
        print(f'MSLE: {msle:10.2f}')

    return (rmse, mae, msle)
```

1.9.1 Dummy

```
[24]: dmr_score = get_scores(regressor=dmr,
                             X_test=X_test,
                             y_true=y_test)
```

RMSE: 220344.24

MAE: 163404.90

MSLE: 0.14

1.9.2 Multilayer Perceptron

```
[25]: mlp_score = get_scores(regressor=mlp,
                             X_test=X_test,
                             y_true=y_test)
```

```
RMSE: 120134.95
MAE:   78720.34
MSLE:    0.04
```

1.9.3 Gradient Boosting Regressor

```
[26]: gbr_score = get_scores(regressor=gbr,
                             X_test=X_test,
                             y_true=y_test)
```

```
RMSE: 115726.26
MAE:   75875.53
MSLE:    0.03
```

1.9.4 Voting Regressor

```
[27]: vote_score = get_scores(regressor=vote,
                              X_test=X_test,
                              y_true=y_test)
```

```
RMSE: 112840.52
MAE:   73245.76
MSLE:    0.03
```

1.9.5 Comparison

We are happy to see that the `VotingRegressor` outperforms the `DummyRegressor` model as well the `GradientBoostingRegressor` and the `MLPRegressor`.

```
[28]: scores = [dmr_score,
                mlp_score,
                gbr_score,
                vote_score]

scores = pd.DataFrame(scores,
                      index=['DMR', 'MLP', 'GBR', 'VOTE'],
                      columns=['RMSE', 'MAE', 'MSLE'])

scores
```

```
[28]:           RMSE      MAE  MSLE
DMR  220,344.24 163,404.90  0.14
```

MLP	120,134.95	78,720.34	0.04
GBR	115,726.26	75,875.53	0.03
VOTE	112,840.52	73,245.76	0.03

1.10 Visualizations

We produce a couple of plots to visually inspect the performance of our model. We use the test data set with the predicted `Amount` to produce the plots.

```
[29]: cols = ['Amount', 'Predicted Amount',
              'District', 'Area', 'Total Rooms']

X_pred = X_test.copy()
X_pred.loc[:, 'Amount'] = y
X_pred.loc[:, 'Predicted Amount'] = vote.predict(X_test)
X_pred = X_pred.loc[:, cols]

X_pred.head()
```

```
[29]:
```


	Amount	Predicted Amount	District	Area	Total Rooms
2776	782000	849,866.16	stare miasto	59	3
3142	417000	401,471.17	pradnik czerwony	47	4
640	239000	281,984.44	pradnik czerwony	24	2
171	230000	347,420.45	podgorze	35	3
1209	889000	879,570.56	bronowice	125	6

On our first visual it can be seen that there exists a fairly linear relationship between the `Predicted Amount` and the `Area` of the property.

```
[30]: plt.scatter(X_pred['Area'], X_pred['Predicted Amount'], s=2)
plt.xlabel('Area')
plt.ylabel('Predicted Amount (PLN)')

ax = plt.gca()
ax.xaxis.set_major_locator(MaxNLocator(integer=True))

plt.tight_layout()
plt.savefig('img/area_vs_amount.png')
plt.show()
```




On the second visual it can be seen, as expected the more Total Rooms in a Property the more it should cost.

```
[31]: plt.scatter(X_pred['Total Rooms'], X_pred['Predicted Amount'], s=2)
plt.xlabel('Total Rooms')
plt.ylabel('Predicted Amount (PLN)')

ax = plt.gca()
ax.xaxis.set_major_locator(MaxNLocator(integer=True))

plt.tight_layout()
plt.show()
```



Next we want to check if the model distinguishes between districts. We group the data by `District` and calculate the mean of the predictions with the group. We produce a bar chart sorted from highest average to lowest. Clearly the model distinguishes between district that are near the city center (`stare miasto`, `zwierzyniec`) and those further away (`łagiewniki`, `bieżanów`).

```
[32]: width = 1600
height = width/2
dpi = 200


X_grp = X_pred[['District', 'Predicted Amount']]
X_grp = X_grp.groupby('District', as_index=False).mean()
X_grp = X_grp.sort_values('Predicted Amount', ascending=False)

plt.figure(figsize=(width/dpi, height/dpi))

plt.bar(X_grp['District'], X_grp['Predicted Amount'] / 1000)

plt.ylabel('Mean predicted amount (k PLN)')
plt.ylim(X_grp['Predicted Amount'].min() * 0.67 / 1000, None)
plt.xticks(rotation=90)

plt.tight_layout()
plt.savefig('img/district_vs_avg_amount.png')
plt.show()
```



1.11 Getting predictions

Next we would like to see how the model handles sets of arbitrary parameters. We write a function to transform inputs to desired format and obtain prediction from the model.

```
[33]: def get_pred(district,
                seller,
                area,
                rooms,
                bathrooms,
                parking,
                garden,
                balcony,
                terrace,
                floor,
                new,
                estate,
                townhouse,
                apartment,
                land,
                studio):

    columns = ['District',
               'Seller',
               'Area',
               'Rooms',
               'Bathrooms',
               'Parking',
               'Garden',
```

```

        'Balcony',
        'Terrace',
        'Floor',
        'New',
        'Estate',
        'Townhouse',
        'Apartment',
        'Land',
        'Studio',
        'Log Area',
        'Bool Sum',
        'Area to Bool Sum',
        'Rooms to Bool Sum',
        'Rooms to Bathrooms',
        'Total Rooms',
        'Area to Rooms',
        'Area to Bathrooms',
        'Area to Total Rooms']

log_area = np.log(area)

all_bools = [garden,
             balcony,
             terrace,
             floor,
             new,
             estate,
             townhouse,
             apartment,
             land,
             studio]

bool_sum = sum(all_bools)
area_to_bool_sum = area / (bool_sum + 1)
rooms_to_bool_sum = rooms / (bool_sum + 1)
rooms_to_bathrooms = rooms / bathrooms
total_rooms = rooms + bathrooms
area_to_rooms = area / total_rooms
area_to_bathrooms = area / bathrooms
area_to_total_rooms = area / total_rooms

x = [district,
     seller,
     area,
     rooms,
     bathrooms,
     parking,
```

```

        garden,
        balcony,
        terrace,
        floor,
        new,
        estate,
        townhouse,
        apartment,
        land,
        studio,
        log_area,
        bool_sum,
        area_to_bool_sum,
        rooms_to_bool_sum,
        rooms_to_bathrooms,
        total_rooms,
        area_to_rooms,
        area_to_bathrooms,
        area_to_total_rooms]

x = pd.DataFrame([x], columns=columns)
x = float(vote.predict(x))

return int(round(x, -3))

```

We create lists of inputs for the model to predict.

```

[34]: areas = range(30, 120, 5)
rooms = range(1, 5)
districts = ['stare miasto',
             'bronowice',
             'krowodrza',
             'borek falecki']

```

Next we loop over lists of possible Area's and Room's and plot the outputs. First we check how the model reacts to different districts.

```

[35]: plt.figure()

for d in districts:
    value = list()
    for a in areas:
        pred = get_pred(district=d,
                        seller='realtor',
                        area=a,
                        rooms=2,
                        bathrooms=1,
                        parking='street',

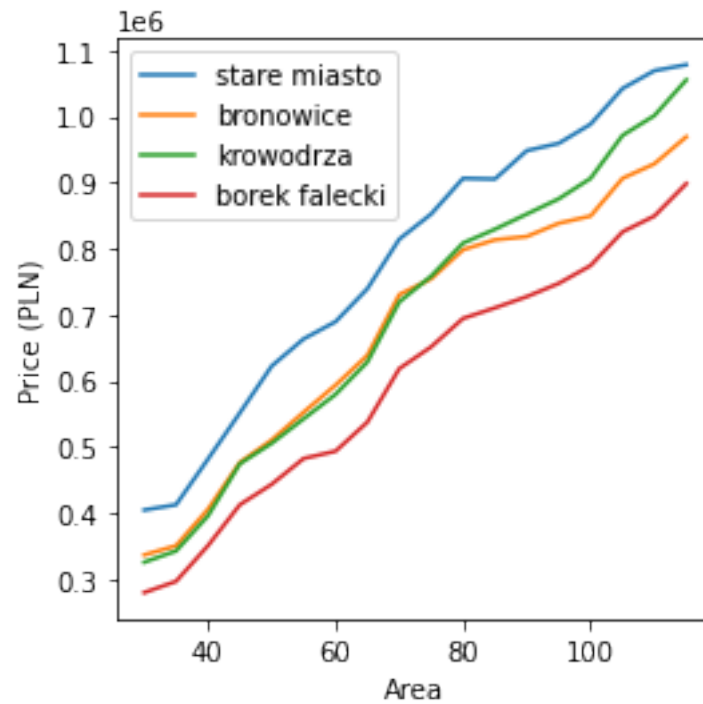
```

```

        garden=False,
        balcony=False,
        terrace=False,
        floor=False,
        new=True,
        estate=False,
        townhouse=True,
        apartment=False,
        land=False,
        studio=True)
    value.append(pred)
plt.plot(areas, value, label=d)

plt.ylabel('Price (PLN)')
plt.xlabel('Area')
plt.legend(loc='best')
plt.savefig('img/area_vs_amount_by_district')
plt.show()

```



We do the same for different amounts of Room's.

```

[36]: plt.figure()

      for r in rooms:

```




```

value = list()
for a in areas:
    pred = get_pred(district='stare miasto',
                    seller='owner',
                    area=a,
                    rooms=r,
                    bathrooms=1,
                    parking='street',
                    garden=False,
                    balcony=True,
                    terrace=False,
                    floor=False,
                    new=True,
                    estate=False,
                    townhouse=True,
                    apartment=False,
                    land=False,
                    studio=True)
    value.append(pred)
plt.plot(areas, value, label=f'{r} rooms')

plt.title('Stare Miasto')
plt.ylabel('Price (PLN)')
plt.xlabel('Area')
plt.legend(loc='best')
plt.savefig('img/area_vs_amount_by_rooms')
plt.show()

```



1.12 Final training

The last step is to fit the model to the entire dataset and save it for later use.

```
[37]: start = datetime.now()

gbr.fit(X, y)
joblib.dump(gbr, f'../flats-model/gbr.joblib')

mlp.fit(X, y)
joblib.dump(mlp, f'../flats-model/mlp.joblib')

vote.fit(X, y)
joblib.dump(vote, f'../flats-model/vote.joblib')

end = datetime.now()

duration = (end - start).seconds

print(f'Full training took {int(duration)} seconds.')
```

Full training took 80 seconds.

```
[38]: # Copy files to portfolio  
# fromDirectory = '.'  
# toDirectory = '/home/dev/Github/data-science-portfolio/flats-in-cracow'  
# copy_tree(fromDirectory, toDirectory)
```